


# Parking Occupancy Detection



# - What we will cover today — — — — —

---



## **Introduction**

Outlining the objective of project



## **Components**

Needed to complete the project



## **Models**

State-of-the-Art to Basic Classification Algorithms



## **Implementation**

What we worked on throughout the project



## **Evaluation**

Comparing the results of our models/C.A prediction




## **Conclusion**

What we've learned and future improvements




Introduction


# Introduction of Project



- Detect whether parking spot is occupied by a vehicle/object or not
- Segment parking spaces from video data
- Real-time monitoring/data on occupancy count
- Can be applied to parking spaces in shopping centres, garages, airports etc.



Components






# — Video of Car Park — — — — —




# Mask of Car Park




# Dataset

**Empty**



**Not Empty**






# Technology Used — — — — —



**OpenCV**

Image/Video Processing



**Matplotlib**

Visual Representation of  
Data



**TensorFlow**

**TensorFlow**


Building, training and  
testing deep learning  
models

# Technology Used



**SKLearn**

Evaluating model  
performance



**NumPy**

**Numpy**

Handling numerical  
operations



**Python**

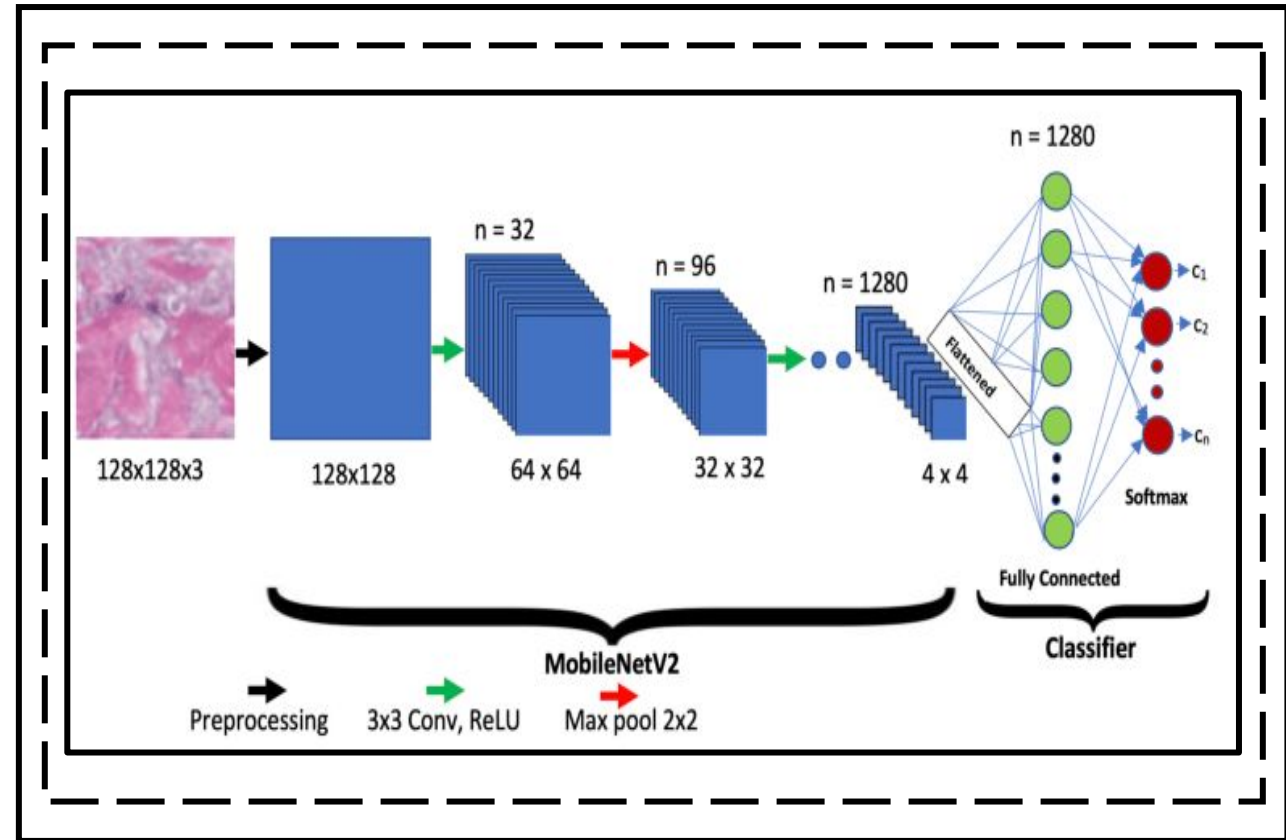
Programming language to  
complete project

3

Models

# MobileNetV2


- Recommended image target size = 224x224 (not optimal for our data set)
- Image size = 32x75 (min 32 required)



# MobileNetV2

```
153/153 ██████████ 6s 28ms/step - accuracy: 0.9045 - loss: 0.2554 - val_accuracy: 0.9458 - val_loss: 0.1669
Epoch 2/10
153/153 ██████████ 4s 26ms/step - accuracy: 0.9952 - loss: 0.0286 - val_accuracy: 0.9606 - val_loss: 0.1474
Epoch 3/10
153/153 ██████████ 4s 27ms/step - accuracy: 0.9971 - loss: 0.0166 - val_accuracy: 0.9680 - val_loss: 0.1155
Epoch 4/10
153/153 ██████████ 4s 28ms/step - accuracy: 0.9997 - loss: 0.0095 - val_accuracy: 0.9672 - val_loss: 0.1036
Epoch 5/10
153/153 ██████████ 4s 27ms/step - accuracy: 0.9999 - loss: 0.0061 - val_accuracy: 0.9721 - val_loss: 0.0968
Epoch 6/10
153/153 ██████████ 4s 27ms/step - accuracy: 1.0000 - loss: 0.0050 - val_accuracy: 0.9713 - val_loss: 0.0979
Epoch 7/10
153/153 ██████████ 4s 28ms/step - accuracy: 1.0000 - loss: 0.0034 - val_accuracy: 0.9672 - val_loss: 0.0966
Epoch 8/10
153/153 ██████████ 4s 28ms/step - accuracy: 1.0000 - loss: 0.0028 - val_accuracy: 0.9680 - val_loss: 0.0969
Epoch 9/10
153/153 ██████████ 4s 28ms/step - accuracy: 1.0000 - loss: 0.0025 - val_accuracy: 0.9680 - val_loss: 0.0936
Epoch 10/10
153/153 ██████████ 4s 28ms/step - accuracy: 1.0000 - loss: 0.0020 - val_accuracy: 0.9680 - val_loss: 0.0938
39/39 ██████████ 1s 22ms/step - accuracy: 0.9730 - loss: 0.0726
Validation Loss: 0.0938146561384201
Validation Accuracy: 0.9679803252220154
```

# Keras Sequential Model



Keras

Image size = 29x68

```
model = Sequential([
    Input(shape=(29, 68, 3)),
    Conv2D(32, (3, 3), activation='relu'),
    MaxPooling2D(2, 2),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D(2, 2),
    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.5),
    Dense(2, activation='softmax')
])
```






internal structure inspired TheDeepHub

# Keras Sequential Model

```
153/153 ————— 3s 16ms/step - accuracy: 0.9033 - loss: 0.2569 - val_accuracy: 0.9655 - val_loss: 0.1186
Epoch 2/10
153/153 ————— 2s 15ms/step - accuracy: 0.9958 - loss: 0.0185 - val_accuracy: 0.9672 - val_loss: 0.1404
Epoch 3/10
153/153 ————— 2s 15ms/step - accuracy: 0.9965 - loss: 0.0086 - val_accuracy: 0.9770 - val_loss: 0.0563
Epoch 4/10
153/153 ————— 2s 15ms/step - accuracy: 1.0000 - loss: 7.9550e-04 - val_accuracy: 0.9844 - val_loss: 0.0369
Epoch 5/10
153/153 ————— 3s 17ms/step - accuracy: 1.0000 - loss: 2.8986e-04 - val_accuracy: 0.9819 - val_loss: 0.0455
Epoch 6/10
153/153 ————— 2s 16ms/step - accuracy: 1.0000 - loss: 1.1871e-04 - val_accuracy: 0.9836 - val_loss: 0.0513
Epoch 7/10
153/153 ————— 2s 16ms/step - accuracy: 1.0000 - loss: 2.0379e-04 - val_accuracy: 0.9836 - val_loss: 0.0387
Epoch 8/10
153/153 ————— 2s 16ms/step - accuracy: 1.0000 - loss: 4.6520e-04 - val_accuracy: 0.9819 - val_loss: 0.0515
Epoch 9/10
153/153 ————— 3s 16ms/step - accuracy: 1.0000 - loss: 7.4225e-05 - val_accuracy: 0.9836 - val_loss: 0.0461
Epoch 10/10
153/153 ————— 3s 16ms/step - accuracy: 1.0000 - loss: 5.2256e-05 - val_accuracy: 0.9811 - val_loss: 0.0542
39/39 ————— 0s 6ms/step - accuracy: 0.9784 - loss: 0.0622
Validation Loss: 0.05424930527806282
Validation Accuracy: 0.9811165928840637
```



# What about HOG Features?



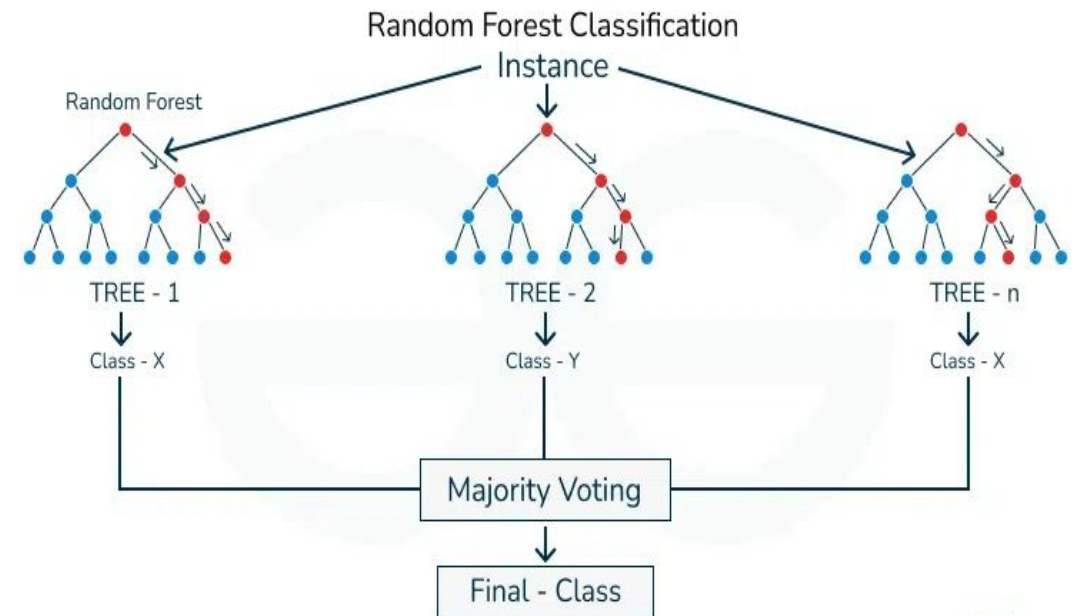
# Random Forest

- Combines the output of multiple decision trees to reach a single result

```
Random Forest Accuracy: 0.9961685823754789
Random Forest Loss: 0.09935789583867653
Random Forest Classification Report:
      precision    recall  f1-score   support

     1       1.00      1.00      1.00     954
     2       1.00      1.00      1.00     873

 accuracy          1.00          1.00          1.00     1827
 macro avg          1.00          1.00          1.00     1827
 weighted avg       1.00          1.00          1.00     1827
```



Random Forest Classifier using Scikit-learn



- Classification tasks: finds the optimal hyperplane that separates different classes in the feature space
- Maximizing the margin between the classes to improve classification accuracy.

```
SVC Accuracy: 1.0
SVC Loss: 0.0005111067423619549
SVC Classification Report:
              precision    recall  f1-score   support

     1           1.00       1.00       1.00        954
     2           1.00       1.00       1.00        873

 accuracy          1.00          1.00          1.00       1827
 macro avg          1.00          1.00          1.00       1827
weighted avg          1.00          1.00          1.00       1827
```

4

Implementation

# — Creating Mask — — — — —

```
def draw_rectangle(event, x, y, flags, param):
    global ix, iy, drawing, img, mask
    if event == cv2.EVENT_LBUTTONDOWN:
        drawing = True
        ix, iy = x, y
    elif event == cv2.EVENT_MOUSEMOVE:
        if drawing:
            mask_copy = mask.copy()
            cv2.rectangle(mask_copy, (ix, iy), (x, y), (255, 255, 255), -1)
            display_img = cv2.addWeighted(img, 0.8, mask_copy, 0.2, 0)
            cv2.imshow("image", display_img)
    elif event == cv2.EVENT_LBUTTONUP:
        drawing = False
        cv2.rectangle(mask, (ix, iy), (x, y), (255, 255, 255), -1)
        display_img = cv2.addWeighted(img, 0.8, mask, 0.2, 0)
        cv2.imshow("image", display_img)
```

```
video_path = r"video.mp4"
cap = cv2.VideoCapture(video_path)

if not cap.isOpened():
    print("Error opening video file")
else:
    ret, frame = cap.read() # Read the first frame
    if not ret:
        print("Failed to read the video")
    else:
        # Resize frame for convenience
        img = cv2.resize(frame, (800, 600)) # Resize for easier handling
        mask = np.zeros_like(img, dtype=np.uint8)

        cv2.namedWindow("image")
        cv2.setMouseCallback("image", draw_rectangle)

        display_img = cv2.addWeighted(img, 0.8, mask, 0.2, 0)
        while True:
            cv2.imshow("image", display_img)
            k = cv2.waitKey(1) & 0xFF
            if k == 27: # ESC key to exit
                break

            cv2.imwrite("mask1.png", mask)
            cv2.destroyAllWindows()

cap.release()
```



# Applying Mask to Video



# Saving Keras Sequential Model



```
PS C:\Users\djibr\Desktop\EMAI\Semester 2 - Sapienza\Computer Vision\Computer Vision> python save_keras_sequential_model.py
Found 4872 images belonging to 2 classes.
Found 1218 images belonging to 2 classes.
Training labels: [1. 1. 1. 1. 1. 0. 1. 1. 1.]
2024-06-20 09:15:57.653512: I tensorflow/core/platform/cpu_feature_guard.cc:182] This TensorFlow binary is optimized to use available CPU instructions
in performance-critical operations.
To enable the following instructions: SSE SSE2 SSE3 SSE4.1 SSE4.2 AVX AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler
flags.
Epoch 1/10
153/153 [=====] - 15s 75ms/step - loss: 0.1259 - accuracy: 0.9557 - val_loss: 0.0905 - val_accuracy: 0.9663
Epoch 2/10
153/153 [=====] - 12s 81ms/step - loss: 0.0152 - accuracy: 0.9955 - val_loss: 0.0848 - val_accuracy: 0.9754
Epoch 3/10
153/153 [=====] - 13s 83ms/step - loss: 0.0108 - accuracy: 0.9967 - val_loss: 0.0443 - val_accuracy: 0.9819
Epoch 4/10
153/153 [=====] - 13s 83ms/step - loss: 0.0085 - accuracy: 0.9973 - val_loss: 0.0376 - val_accuracy: 0.9860
Epoch 5/10
153/153 [=====] - 10s 64ms/step - loss: 0.0023 - accuracy: 0.9994 - val_loss: 0.0456 - val_accuracy: 0.9869
Epoch 6/10
153/153 [=====] - 9s 60ms/step - loss: 0.0041 - accuracy: 0.9984 - val_loss: 0.0117 - val_accuracy: 0.9967
Epoch 7/10
153/153 [=====] - 9s 61ms/step - loss: 0.0038 - accuracy: 0.9990 - val_loss: 0.0467 - val_accuracy: 0.9860
Epoch 8/10
153/153 [=====] - 9s 59ms/step - loss: 3.9094e-04 - accuracy: 1.0000 - val_loss: 0.0114 - val_accuracy: 0.9959
Epoch 9/10
153/153 [=====] - 9s 61ms/step - loss: 1.7091e-04 - accuracy: 1.0000 - val_loss: 0.0175 - val_accuracy: 0.9918
Epoch 10/10
153/153 [=====] - 9s 57ms/step - loss: 0.0015 - accuracy: 0.9994 - val_loss: 0.0074 - val_accuracy: 0.9959
39/39 [=====] - 2s 47ms/step - loss: 0.0074 - accuracy: 0.9959
Validation Loss: 0.007370667532086372
Validation Accuracy: 0.9958949089050293
```



# Bounding Box

```
# Class to represent a bounding box, with the dimensions and border colour initialised
class BoundingBox:
    def __init__(self, x, y, w, h, class_id):
        self.x = x
        self.y = y
        self.w = w
        self.h = h
        self.class_id = class_id

    def draw(self, frame):
        color = (0, 255, 0) if self.class_id == 0 else (0, 0, 255)
        cv2.rectangle(
            frame, (self.x, self.y), (self.x + self.w, self.y + self.h), color, 2
        )
```

```
video_path = r"video.mp4"
cap = cv2.VideoCapture(video_path)

if not cap.isOpened():
    print("Error opening video file")
else:
    ret, frame = cap.read() # Read the first frame
    if not ret:
        print("Failed to read the video")
    else:
        # Resize frame for convenience
        img = cv2.resize(frame, (800, 600)) # Resize for easier handling
        mask = np.zeros_like(img, dtype=np.uint8)

        cv2.namedWindow("image")
        cv2.setMouseCallback("image", draw_rectangle)

        display_img = cv2.addWeighted(img, 0.8, mask, 0.2, 0)
        while True:
            cv2.imshow("image", display_img)
            k = cv2.waitKey(1) & 0xFF
            if k == 27: # ESC key to exit
                break

            cv2.imwrite("mask1.png", mask)
            cv2.destroyAllWindows()

cap.release()
```

# Saving MobileNetV2 Model



```
PS C:\Users\djibr\Desktop\EMAI\Semester 2 - Sapienza\Computer Vision\Computer Vision> python save_mobilenet2_model.py
Found 4872 images belonging to 2 classes.
Found 1218 images belonging to 2 classes.
Training labels: [[0. 1.]
[0. 1.]
[0. 1.]
[0. 1.]
[1. 0.]
[0. 1.]
[0. 1.]
[0. 1.]]
WARNING:tensorflow:`input_shape` is undefined or non-square, or `rows` is not in [96, 128, 160, 192, 224]. Weights for input shape (224, 224) will be
loaded as the default.
2024-06-20 09:20:47.969910: I tensorflow/core/platform/cpu_feature_guard.cc:182] This TensorFlow binary is optimized to use available CPU instructions
in performance-critical operations.
To enable the following instructions: SSE SSE2 SSE3 SSE4.1 SSE4.2 AVX AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler
flags.
Epoch 1/10
153/153 [=====] - 16s 75ms/step - loss: 0.1152 - accuracy: 0.9641 - val_loss: 0.1381 - val_accuracy: 0.9532
Epoch 2/10
153/153 [=====] - 11s 72ms/step - loss: 0.0252 - accuracy: 0.9959 - val_loss: 0.1126 - val_accuracy: 0.9622
Epoch 3/10
153/153 [=====] - 10s 61ms/step - loss: 0.0139 - accuracy: 0.9977 - val_loss: 0.0885 - val_accuracy: 0.9778
Epoch 4/10
153/153 [=====] - 14s 92ms/step - loss: 0.0093 - accuracy: 0.9990 - val_loss: 0.0956 - val_accuracy: 0.9639
Epoch 5/10
153/153 [=====] - 12s 79ms/step - loss: 0.0069 - accuracy: 0.9992 - val_loss: 0.1016 - val_accuracy: 0.9631
Epoch 6/10
153/153 [=====] - 10s 62ms/step - loss: 0.0045 - accuracy: 1.0000 - val_loss: 0.0817 - val_accuracy: 0.9787
Epoch 7/10
153/153 [=====] - 12s 79ms/step - loss: 0.0038 - accuracy: 1.0000 - val_loss: 0.0834 - val_accuracy: 0.9745
Epoch 8/10
153/153 [=====] - 12s 75ms/step - loss: 0.0029 - accuracy: 1.0000 - val_loss: 0.0867 - val_accuracy: 0.9713
Epoch 9/10
153/153 [=====] - 11s 74ms/step - loss: 0.0024 - accuracy: 1.0000 - val_loss: 0.0899 - val_accuracy: 0.9663
Epoch 10/10
153/153 [=====] - 11s 74ms/step - loss: 0.0020 - accuracy: 1.0000 - val_loss: 0.0809 - val_accuracy: 0.9803
39/39 [=====] - 2s 59ms/step - loss: 0.0809 - accuracy: 0.9803
Validation Loss: 0.08088847249746323
Validation Accuracy: 0.9802955389022827
```

# — Resized Training Images — — — — —




## Keras Sequential

```
# Function to resize the desired area for the bounding box to the target size that is specific to the training model used
def preprocess_for_prediction(desired_area, target_size=(68, 29)):
    desired_area_resized = cv2.resize(desired_area, target_size)
    desired_area_normalized = desired_area_resized / 255.0
    desired_area_expanded = np.expand_dims(desired_area_normalized, axis=0)
    return desired_area_expanded
```

## MobileNetV2

```
# Function to resize the desired area for the bounding box to the target size that is specific to the training model used
def preprocess_for_prediction(desired_area, target_size=(75, 32)):
    desired_area_resized = cv2.resize(desired_area, target_size)
    desired_area_normalized = desired_area_resized / 255.0
    desired_area_expanded = np.expand_dims(desired_area_normalized, axis=0)
    return desired_area_expanded
```

# Predicting Bounding Box



```
# Function to draw the bounding boxes on an input frame and predict the occupancy of each box
def draw_bounding_boxes_and_predict(frame, mask, model, bounding_boxes):
    # Turn frame into black and white and mask the different regions/features of the frame
    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
    segmented = cv2.bitwise_and(gray, gray, mask=mask)
    contours, _ = cv2.findContours(
        segmented, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE
    )




    # Finding contours of parking space and add the bounding box to it
    desired_areas = []
    new_bounding_boxes = []
    for contour in contours:
        x, y, w, h = cv2.boundingRect(contour)
        desired_area = frame[y : y + h, x : x + w]
        desired_area_preprocessed = preprocess_for_prediction(desired_area)
        desired_areas.append(desired_area_preprocessed)
        new_bounding_boxes.append(BoundingBox(x, y, w, h, 0)) # Temporary class_id=0

    # Predict occupancy on parking space using trained model
    if desired_areas:
        desired_areas_batch = np.vstack(desired_areas)
        predictions = model.predict(desired_areas_batch)
        predicted_classes = np.argmax(predictions, axis=1)

        for i, bounding_box in enumerate(new_bounding_boxes):
            bounding_box.class_id = predicted_classes[i]



    # Update and draw bounding boxes
    bounding_boxes.clear()
    bounding_boxes.extend(new_bounding_boxes)
    for bounding_box in bounding_boxes:
        bounding_box.draw(frame)

    return frame
```





# — Non-Occupied Parking Spaces — — — — —





```
# Function to display number of non-occupied parking spaces
def draw_empty_spots_counter(frame, bounding_boxes):
    total_spots = len(bounding_boxes)
    empty_spots = sum(bbox.class_id == 0 for bbox in bounding_boxes)
    text = f"Empty Spots: {empty_spots}/{total_spots}"

    # Counter Dimintions
    box_x, box_y = 10, 10
    box_w, box_h = 200, 50

    # Draw counter box and apply text
    cv2.rectangle(
        frame, (box_x, box_y), (box_x + box_w, box_y + box_h), (255, 255, 255), -1
    )
    cv2.putText(
        frame,
        text,
        (box_x + 10, box_y + 30),
        cv2.FONT_HERSHEY_SIMPLEX,
        0.8,
        (0, 0, 0),
        2,
    )

    return frame
```




# SVC/Random Forest

```
# Function to load images and labels from the dataset folder
def load_images(folder, target_size=(68, 29)):
    images = []
    labels = []
    for label, subfolder in enumerate(os.listdir(folder)):
        subfolder_path = os.path.join(folder, subfolder)
        if os.path.isdir(subfolder_path):
            for filename in os.listdir(subfolder_path):
                img_path = os.path.join(subfolder_path, filename)
                if filename.endswith((".png", ".jpg", ".jpeg", ".bmp", ".gif")):
                    img = cv2.imread(img_path, cv2.IMREAD_GRAYSCALE)

                    # Resize image to correct target size
                    if img is not None:
                        img = cv2.resize(img, target_size)
                        images.append(img)
                        labels.append(label)
    return images, labels
```

```
# Load dataset
dataset_folder = "dataset"
images, labels = load_images(dataset_folder)
```


## Load Images/Labels



## Extract Hog Feature

```
# Function to extract HOG features from the images obtained in the dataset
def extract_hog_features(images):
    hog_features = []
    for image in images:
        feature = hog(
            image, pixels_per_cell=(8, 8), cells_per_block=(2, 2), visualize=False
        )
        hog_features.append(feature)
    return np.array(hog_features)
```

```
# Function call to extract HOG features
hog_features = extract_hog_features(images)
```



# Train and Predict

## SVC

```
#####  
# SVC with GridSearchCV #  
#####  
  
# Train the model  
svc_params = {"kernel": ("linear", "rbf"), "C": [1, 10]}  
svc = SVC(probability=True)  
svc_clf = GridSearchCV(svc, svc_params)  
svc_clf.fit(X_train, y_train)  
  
# Predict and display the evaluation of the model  
svc_best = svc_clf.best_estimator_  
svc_predictions = svc_best.predict(X_test)  
svc_probabilities = svc_best.predict_proba(X_test)  
svc_accuracy = accuracy_score(y_test, svc_predictions)  
svc_loss = log_loss(y_test, svc_probabilities)  
print("SVC Accuracy:", svc_accuracy)  
print("SVC Loss:", svc_loss)  
print("SVC Classification Report:\n", classification_report(y_test, svc_predictions))
```

```
# Splitting training and testing data  
X_train, X_test, y_train, y_test = train_test_split(  
    hog_features, labels, test_size=0.3, random_state=42  
)
```

## Random Forest

```
#####  
# Random Forest with GridSearchCV #  
#####  
  
# Train the model  
rf_params = {"n_estimators": [50, 100, 200], "max_depth": [None, 10, 20, 30]}  
rf = RandomForestClassifier()  
rf_clf = GridSearchCV(rf, rf_params)  
rf_clf.fit(X_train, y_train)  
  
# Predict and display the evaluation of the model  
rf_best = rf_clf.best_estimator_  
rf_predictions = rf_best.predict(X_test)  
rf_probabilities = rf_best.predict_proba(X_test)  
rf_accuracy = accuracy_score(y_test, rf_predictions)  
rf_loss = log_loss(y_test, rf_probabilities)  
print("Random Forest Accuracy:", rf_accuracy)  
print("Random Forest Loss:", rf_loss)  
print(  
    "Random Forest Classification Report:\n",  
    classification_report(y_test, rf_predictions),  
)
```





Evaluation



# Keras Sequential

Validation Loss: 0.007370667532086372  
Validation Accuracy: 0.9958949089050293



# MobileNetV2

Validation Loss: 0.08088847249746323  
Validation Accuracy: 0.9802955389022827



# SVC

SVC Accuracy: 1.0  
SVC Loss: 0.0005111067423619549  
SVC Classification Report:

	precision	recall	f1-score	support
1	1.00	1.00	1.00	954
2	1.00	1.00	1.00	873
accuracy			1.00	1827
macro avg	1.00	1.00	1.00	1827
weighted avg	1.00	1.00	1.00	1827



# Random Forest

Random Forest Accuracy: 0.9961685823754789  
Random Forest Loss: 0.09935789583867653  
Random Forest Classification Report:

	precision	recall	f1-score	support
1	1.00	1.00	1.00	954
2	1.00	1.00	1.00	873
accuracy			1.00	1827
macro avg	1.00	1.00	1.00	1827
weighted avg	1.00	1.00	1.00	1827

# Demonstration of Project

---

Live Demonstration done in class

See GitHub Repository to run the  
project



Conclusion

# — What we've learned — — — — —

---

1


Basic Classification  
Algorithms **outperformed**  
Convuluted models based  
on project **objectives**

2

**Quality** and  
**pre-processing** of the  
dataset/labels is  
important

3

Future improvements  
(**Use of IoT devices,**  
**real-time camera feed**)



— — — — — Thank You — — — — —

— — — — — Any Questions — — — — —

